Bayesian Classifiers

Consider the following training data:

x	y
1	+
2	−
1	−
1	−

What should be the classification of a test example if $x = 1$? We can give a probabilistic argument that the best answer is "−". Estimating probabilities from the training data we have:

$$\mathsf{prob}(y = + \mid x = 1) = \frac{1}{3}, \qquad \mathsf{prob}(y = - \mid x = 1) = \frac{2}{3} \tag{1}$$

The classification $y = -$ is called *optimal Bayes*. The general case can be expressed as follows:

$$h = \arg\max_h \ \mathsf{prob}(h \mid D)$$

The hypothesis h determined by this rule is called *the Optimal Bayes* hypothesis. In the example above D is $x = 1$, and there are two hypotheses: h_+ is $y = +$, and h_- is $y = -$. The Optimal Bayes hypothesis is h_-.
Using the Bayes theorem we have:

$$h = \arg\max_h \ \mathsf{prob}(h \mid D) = \arg\max_h \ \frac{\mathsf{prob}(D \mid h)\ \mathsf{prob}(h)}{\mathsf{prob}(D)} = \arg\max_h \ \mathsf{prob}(D \mid h)\ \mathsf{prob}(h)$$

The rule on the right hand side can be used when both $\mathsf{prob}(D \mid h)$ and $\mathsf{prob}(h)$ are known. The probability $\mathsf{prob}(h)$ is called the *a-priori* probability of h. The hypothesis h determined by the above rule is called *a maximum a-posteriori*, or a MAP hypothesis. When only $\mathsf{prob}(D \mid h)$ is known we can use the rule:

$$h = \arg\max_h \ \mathsf{prob}(D \mid h)$$

The hypothesis h determined by the above rule is called *a maximum likelihood*, or a ML hypothesis. To summarize, there are three rules that use different probabilities to determine the right hypothesis:

$h_1 = \arg\max_h \ \mathsf{prob}(h \mid D)$	h_1 is the **Optimal Bayes** hypothesis
$h_2 = \arg\max_h \ \mathsf{prob}(D \mid h) \cdot \mathsf{prob}(h)$	h_2 is the **MAP (maximum *a-posteriori*)** hypothesis
$h_3 = \arg\max_h \ \mathsf{prob}(D \mid h)$	h_3 is the **ML (maximum likelihood)** hypothesis

In the table above both h_1 and h_2 are optimal. If there is a unique optimal hypothesis then $h_1 = h2$. The hypothesis h_3 is not optimal.

From the probabilities in (1) we see that $y = -$ is the Optimal Bayes hypothesis. To compute the MAP hypothesis we need the following probabilities, computed with D standing for $x = 1$.

	$\mathsf{prob}(D \mid h)$	$\mathsf{prob}(h)$
$h_+ :\ y = +$	1	1/4
$h_- :\ y = -$	2/3	3/4

Since $1 \cdot 1/4 < 2/3 \cdot 3/4$ the MAP hypothesis is h_-, same as the Optimal Bayes hypothesis. From the above probabilities we also see that the ML hypothesis is h_+.

Naive Bayesian

Typically the probabilities needed for MAP/ML are not available. To see this, consider the following simple extension of the example, where the feature vector has 3 values instead of 1.

x_1	x_2	x_3	y
1	1	1	+
2	2	2	−
1	1	2	−
1	2	1	−

What should be the classification if a test example is $x_1 = 1, x_2 = 2, x_3 = 2$?

Since the test example is not part of the training data we cannot compute the probabilities needed fo MAP or ML. The **Naive Bayesian** technique allows us to compute an approximate MAP hypothesis by assuming conditional independence of feature values. Specifically, the assumption is as follows:

$$\mathsf{prob}(x_1 = a_1, x_2 = a_2, x_3 = a_3 \mid h) = \mathsf{prob}(x_1 = a_1 \mid h) \cdot \mathsf{prob}(x_2 = a_2 \mid h) \cdot \mathsf{prob}(x_3 = a_3 \mid h)$$

Therefore, following probabilities must be estimated for the two hypotheses:

	$\mathsf{prob}(x_1 = 1 \mid h)$	$\mathsf{prob}(x_2 = 2 \mid h)$	$\mathsf{prob}(x_3 = 2 \mid h)$	$\mathsf{prob}(h)$
$h_+ : \ y = +$	1	0	0	1/4
$h_- : \ y = -$	2/3	2/3	2/3	3/4

Since $1 \cdot 0 \cdot 0 \cdot 1/4 < 2/3 \cdot 2/3 \cdot 2/3 \cdot 3/4$ the MAP hypothesis is h_-. Since $1 \cdot 0 \cdot 0 < 2/3 \cdot 2/3 \cdot 2/3$ the ML hypothesis is also h_-.

The general Naive Bayesian rule is as follows. Given a test example $x = (x_1 = a_1 \ldots, x_n = a_n)$. Suppose there are k possible labels $v_1, \ldots, v_k$. Estimate the following probabilities from the training data:

$$\begin{aligned}
&\mathsf{prob}(x_1 = a_1 \mid h = v_1), \ \ldots, \ \mathsf{prob}(x_n = a_n \mid h = v_1), \ \mathsf{prob}(h = v_1) \\
&\mathsf{prob}(x_1 = a_1 \mid h = v_2), \ \ldots, \ \mathsf{prob}(x_n = a_n \mid h = v_2), \ \mathsf{prob}(h = v_2) \\
&\vdots \\
&\mathsf{prob}(x_1 = a_1 \mid h = v_k), \ \ldots, \ \mathsf{prob}(x_n = a_n \mid h = v_k), \ \mathsf{prob}(h = v_k)
\end{aligned}$$

Then compute the following k values:

$$\begin{aligned}
q_1 &= \mathsf{prob}(x_1 = a_1 \mid h = v_1) \cdot \ldots \cdot \mathsf{prob}(x_n = a_n \mid h = v_1) \cdot \mathsf{prob}(h = v_1) \\
q_2 &= \mathsf{prob}(x_1 = a_1 \mid h = v_2) \cdot \ldots \cdot \mathsf{prob}(x_n = a_n \mid h = v_2) \cdot \mathsf{prob}(h = v_2) \\
&\vdots \\
q_k &= \mathsf{prob}(x_1 = a_1 \mid h = v_k) \cdot \ldots \cdot \mathsf{prob}(x_n = a_n \mid h = v_k) \cdot \mathsf{prob}(h = v_k)
\end{aligned}$$

The Naive Bayesian decides on the label v_i if $q_i = \max(q_1, \ldots q_k)$.